

Hadoop EC2 Assignment

Brent Smith, CSCI5673 Spring 2011

Dataset

- ▶ TransStats Aviation
- ▶ Provided by the Bureau of Transportation Statistics to the public
- ▶ Older copy of the data is already available via Amazon's public data sets (<http://aws.amazon.com/datasets/2289>)
- ▶ DBIBCoupon table that contains flight data for each quarter.
 - ▶ Extracted all quarters in 2007 and concatenated into single file for the year.
 - ▶ Roughly 6GB.
 - ▶ Copied to 8GB EBS storage volume
- ▶ Format is one line per flight, in CSV form




Processing

- ▶ Goal was to aggregate all the data by the origin, destination and airline.
- ▶ Aggregator.java
 - ▶ Map Task
 - ▶ Fields to group on is passed on cmd line (also #mappers and #reducers)
 - ▶ Input
 - Key is ignored
 - Value is line from CSV file
 - ▶ Output
 - Key = grouped fields (origin, destination, airline)
 - Value = Literal ONE
 - ▶ Reduce Task
 - ▶ Sum list of values for each key (origin, destination, airline)
 - ▶ Output
 - Key = grouped fields (origin, destination, airline)
 - Value= sum of all values for given key (# flights with same origin, destination, airline)



Results

- ▶ Combination of Origin+Destination = Route
- ▶ Total of ~150,000 unique routes+airlines
- ▶ Top 10 most frequent routes plotted
- ▶ Makes it easy to spot “hubs”
 - ▶ Northwest* = MSP (Minneapolis, Minnesota)
 - ▶ American Airlines = DFW (Dallas, Ft. Worth)
 - ▶ Delta = ATL (Atlanta, Georgia)
 - ▶ *Northwest since merged with Delta.



Experimental Setup


- ▶ **Four tests**
 - ▶ Default* mappers, One Reducer
 - ▶ *6GB / 64MB = ~ 92 mappers
 - ▶ 100 mappers, #Reducers = #nodes in cluster (5/10/19)
 - ▶ Default mappers + Combine stage, One Reducer
 - ▶ 100 mappers + Combine stage, #Reducers = #nodes in cluster
- ▶ **Ran these four tests on clusters with size 5, 10 and 19**
 - ▶ Excludes the manager node
 - ▶ Can only run 20 instances on EC2 so only 19 possible workers
- ▶ **Ran on both m1.small and c1.medium EC2 instance types**
- ▶ **First test was run with empty buffer cache**
 - ▶ Subsequent tests (two, three, four) may be affected by OS caching



Performance Analysis (m1.small)

- ▶ m1.small instance type
- ▶ Interesting result: 1 Reducer + Combine stage outperforms many reducers plus combiner
 - ▶ Why? Additional overhead to schedule multiple reducers?
 - ▶ Difference is quite small.

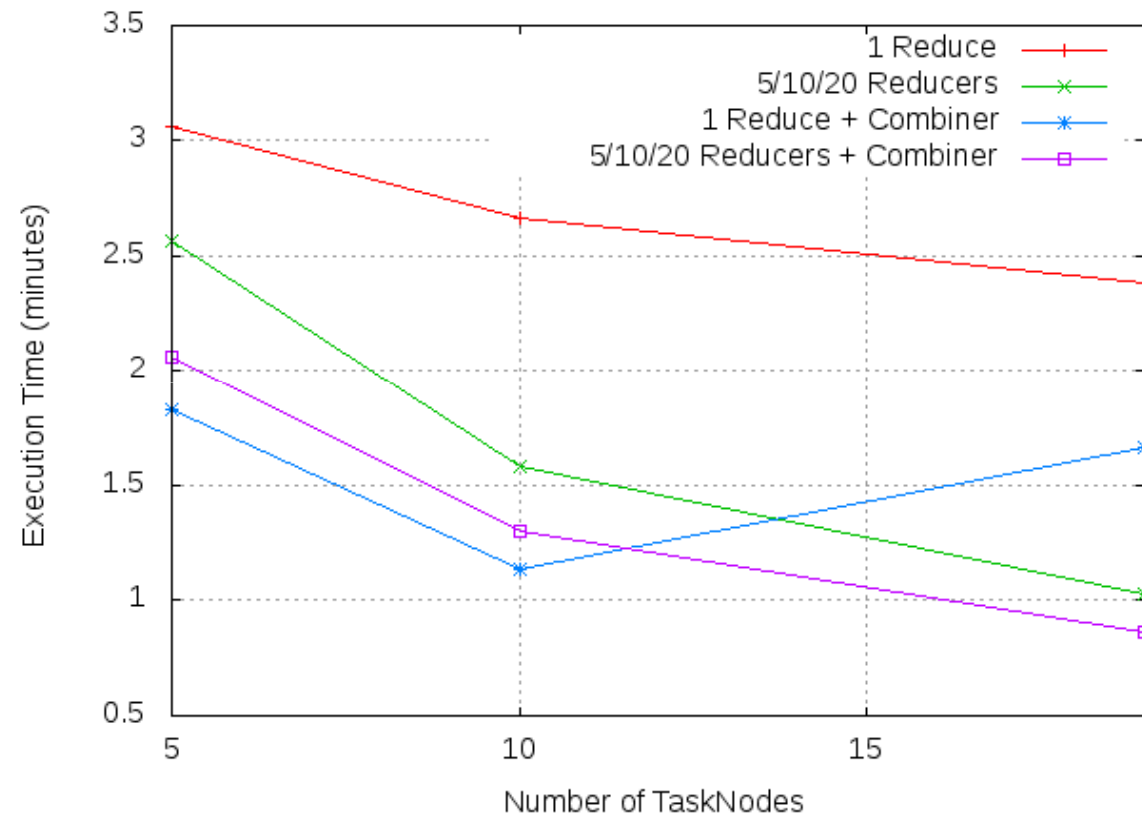
Performance of Aggregation Task on Amazon EC2 Using m1.small Instances



Performance Analysis (c1.medium)

- ▶ c1.medium instance type
- ▶ Interesting result: 1 Reducer + Combine stage performs poorly with 20 nodes
 - ▶ Why?
Bandwidth/data transfer limit? One reducer has to pull all output data from 92 files from each map task?

Performance of Aggregation Task on Amazon EC2 Using c1.medium Instances



Extra Credit

- ▶ Wanted to do something a little more challenging with the data to understand how difficult it is to do something other than simple aggregation.
- ▶ IDEA: Use a graph algorithm to figure out the largest complete subgraph of the graph of all origins and destinations.
 - ▶ Maximal clique problem in graph theory.
 - ▶ Why? Should tell you the largest set of airports for which every airport in the set has a direct route to every other airport in the set.
 - ▶ Is this useful? I don't know. Maybe if you were on the run from the law.



GraphGeneration

- ▶ Implemented as a chained MapReduce job.
 - ▶ Uses the output from the Aggregator MR job as input.
- ▶ GraphGenerator.java
 - ▶ Map Task
 - ▶ Takes a threshold on # of flights. Don't consider routes with less than 365 flights, e.g. less than one flight per day.
 - ▶ Input
 - Key is Origin, Destination, Airline
 - Value is # of flights
 - ▶ Output
 - Key = Origin
 - Value = Destination
 - ▶ Reduce Task
 - ▶ Create a java.util.Set of destinations from list of values
 - ▶ Output
 - Key = Origin
 - Value = distinct set of destinations
 - ▶ Output creates adjacency list for directed graph



Maximal Clique

- ▶ Hard to understand how to do this with MR...
- ▶ But, somebody else has already done it
 - ▶ XRIME: Cloud-Based Large Scale Social Network Analysis
 - ▶ http://ieeexplore.ieee.org/xpl/freeabs_all.jsp?reload=true&arnumber=5557271 (Wei Xue; JuWei Shi; Bo Yang; IBM Res. China, Beijing, China)
- ▶ **XRIME**
 - ▶ Implements many graph algorithms using Hadoop MR
 - ▶ Implements maximal clique
- ▶ Solution: stand on the shoulders of others; e.g., use XRIME.



Chained Maximal Clique

- ▶ Input to Maximal Clique MR Job in XRIME is adjacency list
 - ▶ Corresponds to output from GraphGenerator MR job
- ▶ ChainedMaximalCliqueJob.java
 - ▶ Run Aggregation task
 - ▶ Run GraphGenerator task
 - ▶ Run TextAdjTransformer task
 - ▶ Provided by XRIME
 - ▶ Converts text based adjacency list into binary form
 - ▶ Run MaximalStrongCliqueAlgorithm
 - ▶ Outputs value that is set of all nodes in maximal clique
 - ▶ Use temporary directories to chain output from one MR job to input of another MR job



Maximal Clique Results

- ▶ Many “maximal” cliques of various sizes
 - ▶ Because “maximal” clique is defined as largest set of nodes for which adding another node would make subgraph incomplete
 - ▶ We want the largest “maximal” clique
 - ▶ Use `largest_airport_cliques.py` to parse output and find these.
- ▶ **4 maximal cliques of size 27.**
 - ▶ Complete subgraphs, so # edges = $n(n-1)/2 = 351$ (undirected graph)
 - ▶ Can reach any airport in set from any other airport in set
 - ▶ Useful if you want to increase the work that law enforcement has to do to catch you
- ▶ Four sets are available in `README.txt`.



Questions

- ▶ Still not sure how to do maximal clique with MR...
 - ▶ Source code is available, study it.
- ▶ Maximal clique operates on undirected graphs
 - ▶ Ours was directed.
 - ▶ Assume that XRIME converts directed to undirected only when there are two directed edges ($A \rightarrow B$, and $B \rightarrow A$).
 - ▶ MR sub-jobs in XRIME called “StrongNeighborhoodGenerate”; I think these have something to do with directed \rightarrow undirected conversion
 - ▶ Read XRIME paper, maybe it talks about this?

